


Verifier-Gated Agentic Framework for Credible Cross-Family Model Selection

structured proposals -> hard verification -> non-test evidence -> policy selection -> claim-boundary audit

non-test verifier/evidence feedback refines next proposal round



Cross-family hypothesis space:

Forecasting baselines

Manual mechanistic models

Observation-aware structured search

Discovery ablations

Verification, Evidence Discipline, and Claim-Boundary Auditing

● no-leakage prompt context

● train/validation or rolling evidence only

● post-hoc test metrics after selection

● compact reproducible artifacts

● claim strength matched to evidence

Agreement across families/proposers/evidence modes increases confidence

Disagreement narrows the recommendation instead of forcing a global winner

Rejected: SOTA, FluSight leaderboard, autonomous science, mechanism recovery, provider superiority